

# House Prices: Advanced Regression Techniques

**Predict sales prices and practice feature engineering, RFs, and gradient boosting**

by a1804938

School of Computer Science, The University of Adelaide

Report submitted for **4020\_COMP\_SCI\_7209 Big Data Analysis and Project** at, University of Adelaide towards the Master of Data Science



## Abstract

Testing accuracy of 87.04% was achieved using LGBM Regressor while a testing accuracy of 86.62% was achieved using XGB regressor. Model was trained on data acquired from Ames Housing Dataset compiled by Dean De Cock[1]. It is achieved by intensive Explanatory data analysis by dividing the data into categorical and numerical values, handling of missing and null data, applying encoding on the categorical data and then fitting the data into a model to calculate it's accuracy and root square result.

## Introduction

In this world everyone dreams of owning a house and based on that dream is how the real estate market survives and hence making it the biggest and one of the most important decision an individual have to make in their lifetime. This project is aimed at predicting the housing sale prices in Ames, by using various features and aspects of any individual home. The various aspects and the process of training the data gives any individual insight on making the decision to make a purchase helping them to maximize their gain keeping in mind about their budget. These findings also help real estate agents in improving their sales based on marketing of some key features.



*Main street in downtown Ames, Photo by Tim Kiser*

The Ames Housing Data Set is an alternative to one of the most known and worked on dataset - Boston Housing Data Set. The Ames Housing Dataset was collected by a Professor at Truman State University, Dean De Cock, in 2011. This dataset comprises of residential properties in the town of Ames, ranging from year 2006 to 2010, consisting of 80 features with 2930 samples. This dataset is divided into training and test data sets. Training data set comprises of 1460 observations and testing dataset comprises of 1459 observations.

From the dataset following details can be understood:

- A total of 80 features of which is divided into categorical and numerical data.
- There are total of 46 categorical features ranging from 2 to 28 classes like largest neighborhood etc.
- There are 14 discrete items like kitchens, baths etc.
- There are 20 continuous items like square footage, lot size etc.

Categorical features [2] are the ones which might contain few distinct groups or finite number of categories and not necessarily have a logical order to the data. Categorical data examples are like gender or mode of payment. Gender maybe stored as male = 1 or female = 0, which can't be used to directly perform arithmetic operations. Numerical data features are like quantitative characteristics which are expressed as quantitative values and can perform some arithmetic operations.

Numerical data is further divided into discrete and continuous variables.

Discrete variable is the data that have countable number of values between 2 values. For example, number of defects or issues.

Continuous variable is the data that have an infinite number of values between 2 values which either in the form of numbers or date time. For example, a transaction data and time.

The training dataset includes the sale price, hence making it a supervised regression machine learning task. A supervised task is one where we have access to both target and features. We have to train a model which can do mapping between both features and target. Sale price which is the target is a continuous variable. The model should be accurate i.e. it should be able to predict the sale price nearest to its actual value and the model predictions should be interpretable.

### Step by step methodology that can be followed for the analysis of this dataset:

1. EDA - Exploratory data analysis
2. Data handling – formatting and cleaning
3. Feature engineering
4. Try to compare various machine learning models for performance
5. Perform parameter tuning for the selected model
6. Train models on data set
7. Submit predictions

Data formatting and cleaning take care of:

- Null or missing values
- Wrong datatype
- Wrong data
- Outliers
- Skewed distributions

### Exploratory data analysis

```
In [3]: # Load the dataset
#house = pd.read_csv("/kaggle/input/house-prices-advanced-regression-techniques/train.csv")
#house.head()
df = pd.read_csv("/Users/user/Downloads/all/train.csv")
df.head()
```

```
Out[3]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	Condition2
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	Feedr
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	FR2	Gtl	Veenker	Feedr	Feedr
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	Feedr
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	Corner	Gtl	Crawfor	Norm	Feedr
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	FR2	Gtl	NoRidge	Norm	Feedr

```
In [4]: df.shape
```

```
Out[4]: (1460, 81)
```

We have a total of 81 columns. Target variable is SalePrice that is to be predict. Id column is just an index which is to be ignored. So we have a total 79 input features to predict the output from.

#	Column	Non-Null	Count	Dtype			
---	-----	-----	-----	-----			
0	Id	1460	non-null	int64	28	ExterCond	1460 non-null object
1	MSSubClass	1460	non-null	int64	29	Foundation	1460 non-null object
2	MSZoning	1460	non-null	object	30	BsmtQual	1423 non-null object
3	LotFrontage	1201	non-null	float64	31	BsmtCond	1423 non-null object
4	LotArea	1460	non-null	int64	32	BsmtExposure	1422 non-null object
5	Street	1460	non-null	object	33	BsmtFinType1	1423 non-null object
6	Alley	91	non-null	object	34	BsmtFinSF1	1460 non-null int64
7	LotShape	1460	non-null	object	35	BsmtFinType2	1422 non-null object
8	LandContour	1460	non-null	object	36	BsmtFinSF2	1460 non-null int64
9	Utilities	1460	non-null	object	37	BsmtUnfSF	1460 non-null int64
10	LotConfig	1460	non-null	object	38	TotalBsmtSF	1460 non-null int64
11	LandSlope	1460	non-null	object	39	Heating	1460 non-null object
12	Neighborhood	1460	non-null	object	40	HeatingQC	1460 non-null object
13	Condition1	1460	non-null	object	41	CentralAir	1460 non-null object
14	Condition2	1460	non-null	object	42	Electrical	1459 non-null object
15	BldgType	1460	non-null	object	43	1stFlrSF	1460 non-null int64
16	HouseStyle	1460	non-null	object	44	2ndFlrSF	1460 non-null int64
17	OverallQual	1460	non-null	int64	45	LowQualFinSF	1460 non-null int64
18	OverallCond	1460	non-null	int64	46	GrLivArea	1460 non-null int64
19	YearBuilt	1460	non-null	int64	47	BsmtFullBath	1460 non-null int64
20	YearRemodAdd	1460	non-null	int64	48	BsmtHalfBath	1460 non-null int64
21	RoofStyle	1460	non-null	object	49	FullBath	1460 non-null int64
22	RoofMat1	1460	non-null	object	50	HalfBath	1460 non-null int64
23	Exterior1st	1460	non-null	object	51	BedroomAbvGr	1460 non-null int64
24	Exterior2nd	1460	non-null	object	52	KitchenAbvGr	1460 non-null int64
25	MasVnrType	1452	non-null	object	53	KitchenQual	1460 non-null object
26	MasVnrArea	1452	non-null	float64	54	TotRmsAbvGrd	1460 non-null int64
27	ExterQual	1460	non-null	object	55	Functional	1460 non-null object
					56	Fireplaces	1460 non-null int64


We can see that there's some logical order in the columns to go from general information to detailed ones and describe the neighborhood area and outside area first like Street, alley, LotArea, and then the inside information which goes from basement, heating, floors to rooms etc. with final information about the sale.

There are lots of null values present in the dataset. 4 of these features have missing values more than 50%.

	column_name	percentage
72	PoolQC	99.520548
74	MiscFeature	96.301370
6	Alley	93.767123
73	Fence	80.753425
57	FireplaceQu	47.260274
3	LotFrontage	17.739726
58	GarageType	5.547945
59	GarageYrBlt	5.547945
60	GarageFinish	5.547945
63	GarageQual	5.547945
64	GarageCond	5.547945
32	BsmtExposure	2.602740
35	BsmtFinType2	2.602740
33	BsmtFinType1	2.534247
31	BsmtCond	2.534247
30	BsmtQual	2.534247
26	MasVnrArea	0.547945
25	MasVnrType	0.547945
42	Electrical	0.068493


```
#distribution of saleprice
sns.distplot(df.SalePrice)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x20a0b9f11c0>
```

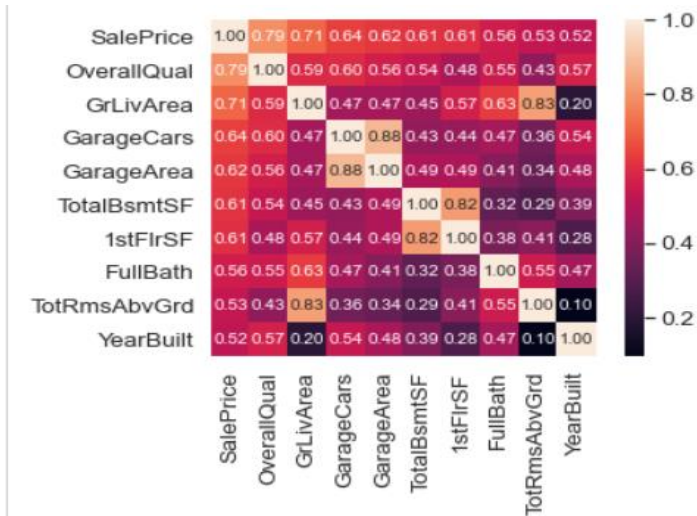


SalePrice seems to be skewed. Skewness is the amount of distortion when compared to symmetrical bell curve or normal distribution measuring the lack of symmetry in data distribution.[3] The values seem they aren't normally distributed.

Correlation of the data can be seen as below:



The correlation of saleprice with 9 other important columns can be seen as below:




Then we can see the scatter plot of these variables with salesprice to have a better understanding of the data and their dependency on salesprice.



Then the Id column is dropped from the dataset as it is not required for the processing of the model.

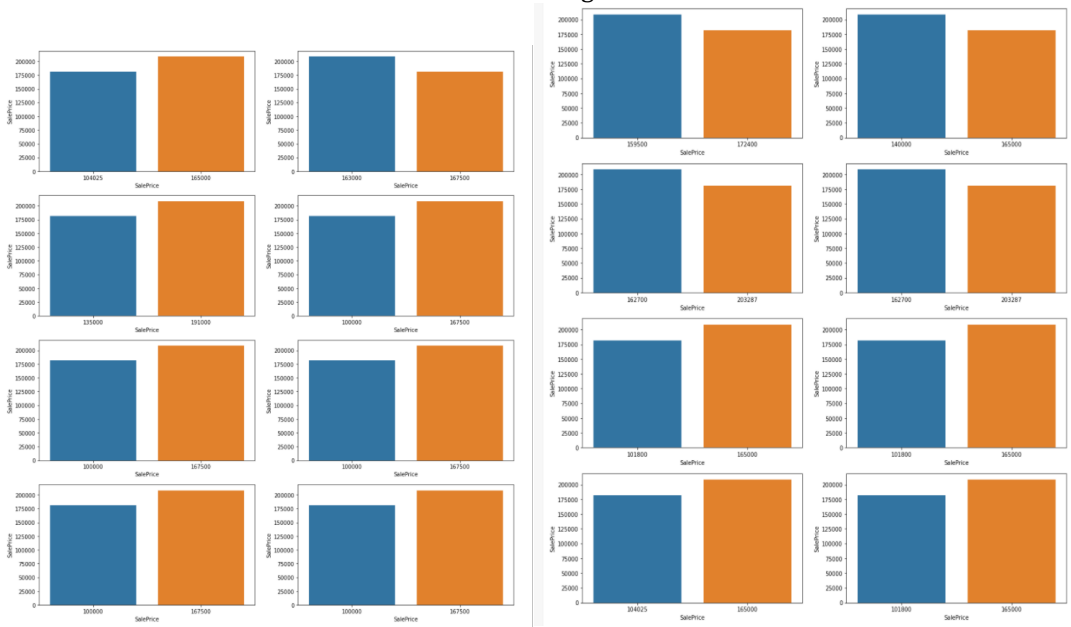
Then null values handling is done. First checked the null values in the columns:



After checking the data dictionary, these are not actually the null values, rather these are the features which are not present in the house.

For example, let check the field Alley, Value "NA: here means house has "No Alley Access"

Then the correlation of these null values with the target variable is checked.



From the above graphs, we can clearly see that the null values have strong relation with the SalePrice, hence we can neither drop the columns with null values, nor we can drop the rows with null values.

To handle these null and missing values below steps are performed:


- 1) All missing values for the categorical columns will be replaced by "None"
- 2) All missing values for the numeric columns will be replaced by median of that field

Then the date variables are examined. The following are the date variables:

GarageYrBlt', 'YearBuilt', 'YearRemodAdd', 'YrSold'

The relation of these variables are examined with the salesprice.





The trend of sale price with the feature "YrSold", it shows a decreasing trend. In real estate this seems to be unreal because as time passes the price is expected to increase which seems to be opposite over here. This is handled by creating a variable "Age" from these variables.

3 new parameters are created:

Age variable:- 'HouseAge' = 'YrSold' - 'YearBuilt'

Age of master after remodeling:- 'RemodAddAge' = 'YrSold' - 'YearRemodAdd'

Creating age of the garage from year built of the garage to the sale of the master

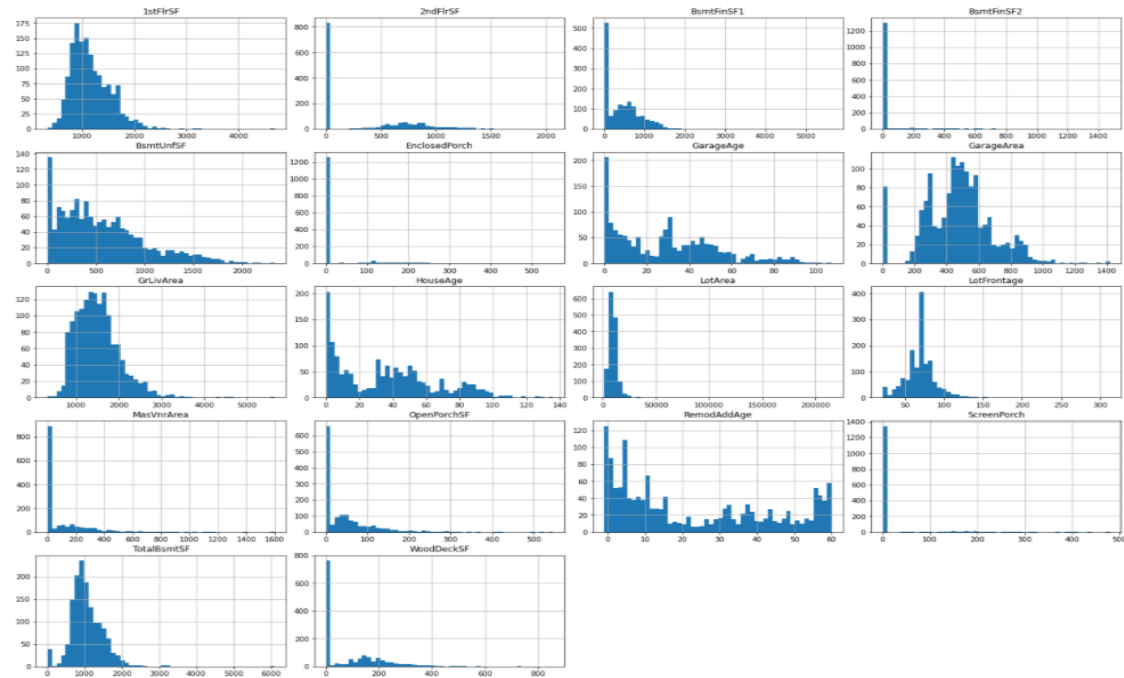
'GarageAge' = 'YrSold' - 'GarageYrBlt'

And then the original variables: "YearBuilt", "YearRemodAdd", "GarageYrBlt" are dropped from the dataset.

Then a separate lists of categorical and numeric columns are created based on the datatypes. Then the numeric features are separated into continuous and discrete numeric features.

Then the variance in the different continuous numeric columns present in the dataset is examined which can be seen in the below scatter plot.

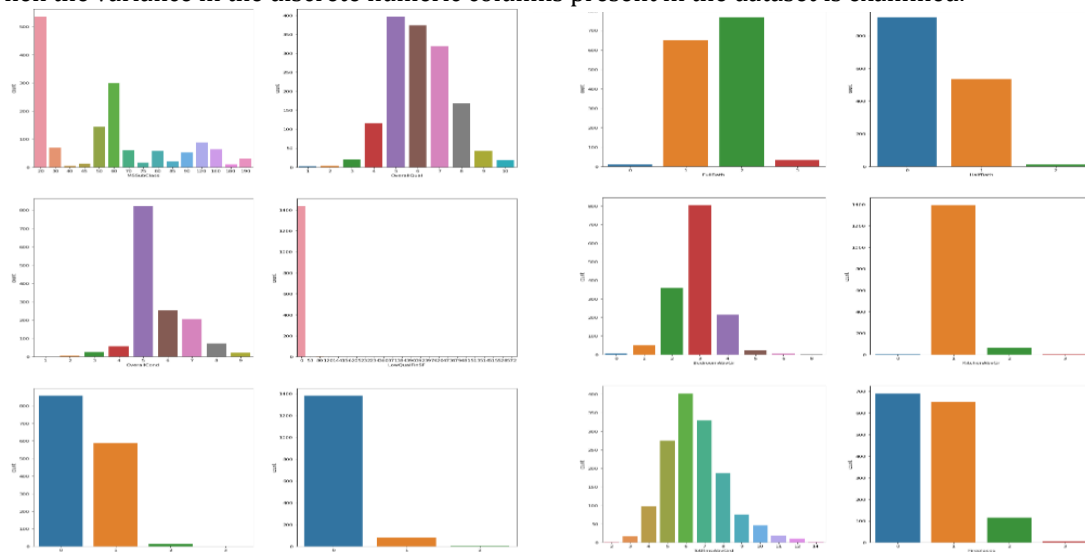


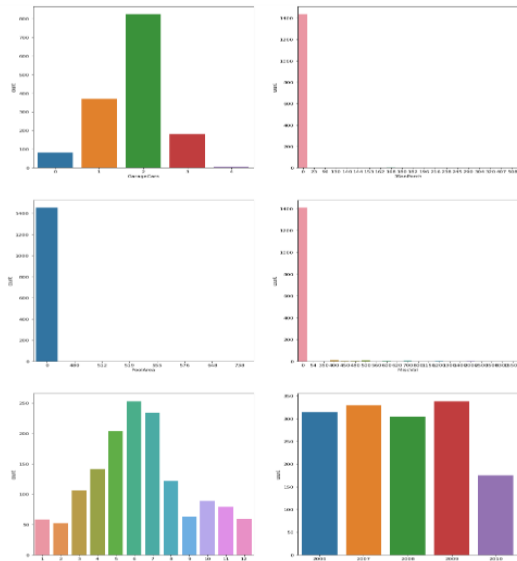


Then the variance in numbers is examined to find out the list of features that doesn't have much effect on the target variables by checking the count of the values of those column values. Following variables seems to have low variance:

MasVnArea, BsmFinSF2, 2ndFlrSF, EnclosedPorch, ScreenPorch.

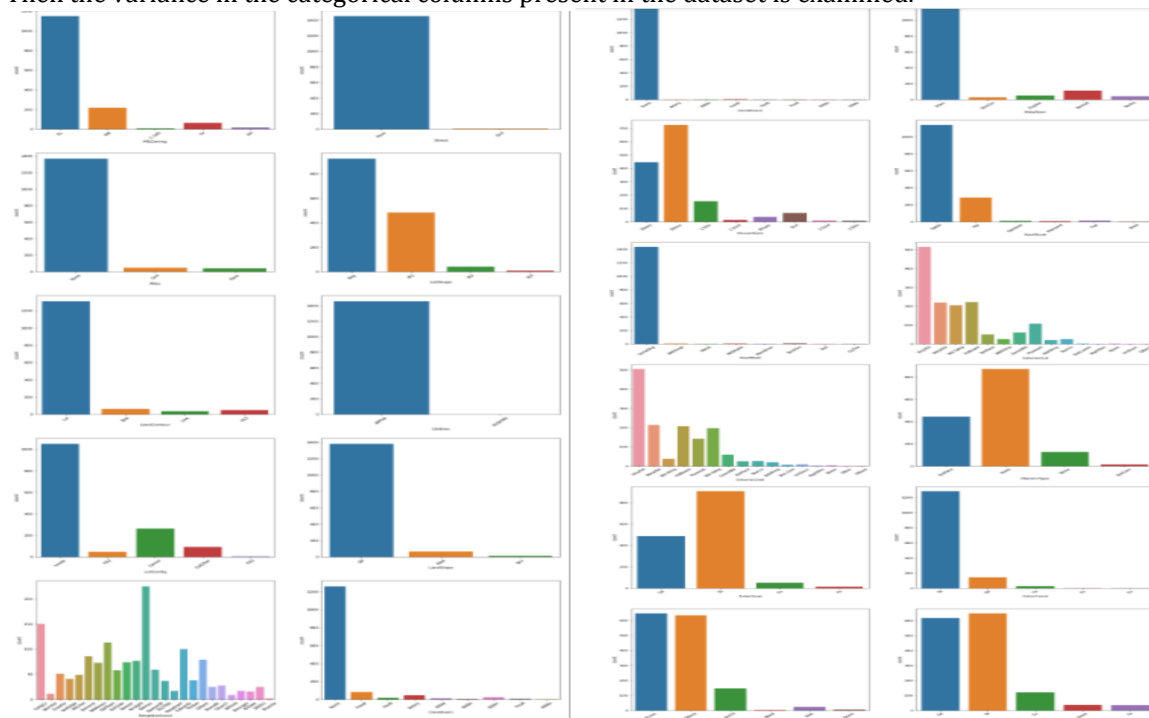
Then the variance in the discrete numeric columns present in the dataset is examined.





The following variables seems to have low variance:  
 LowQualFinSF, BsmtHalfBath, KitchenAbvGr, 3SsnPorch, PoolArea, MiscVal.

Then the variance in the categorical columns present in the dataset is examined.



Following variables seems to have low variance:


MSZoning Street, Alley LandContour, Utilities, LotConfig Condition1 LandSlope Condition2, BldgType RoofStyle, RoofMatl, ExterCond, BsmtCond, BsmtFinType2, Heating, CentralAir, Electrical Functional GarageQual, GarageCond, PavedDrive, PoolQC, Fence, MiscFeature, SaleType, SaleCondition.

All the columns which have low variables as identified are dropped from the dataset.

Then the skewness in saleprice, is handled by taking the log to get normal distribution which becomes like below.

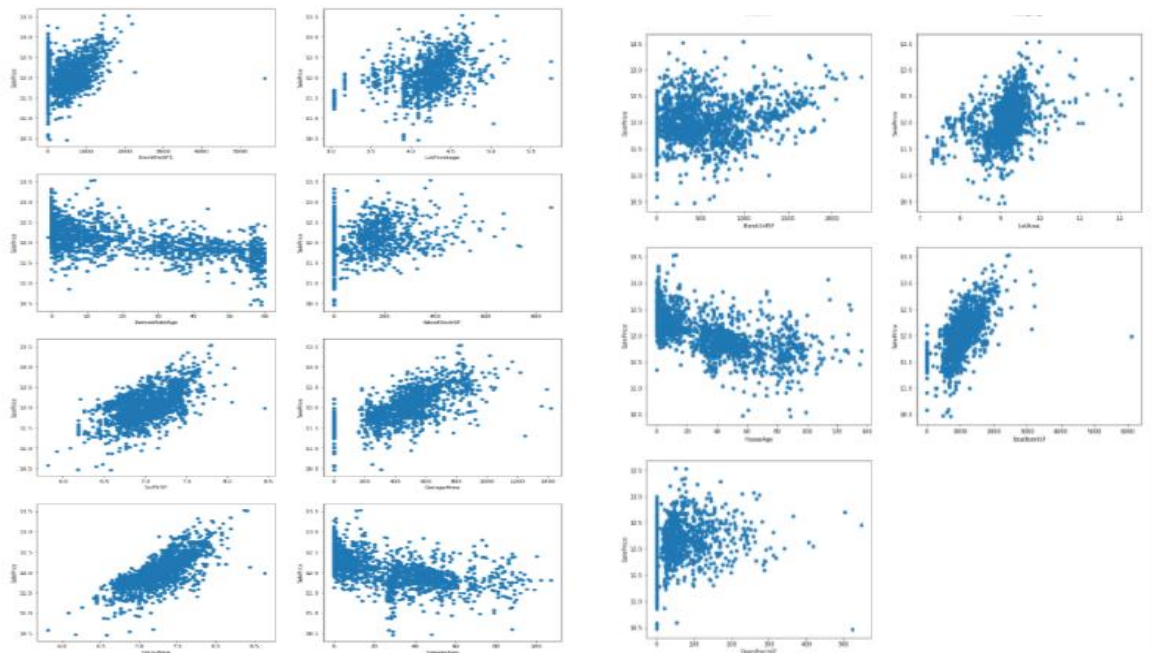
```
sns.distplot(df.SalePrice)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x13e87c97c:
```



Also the log of few numeric variables 'LotFrontage', 'LotArea', '1stFlrSF', 'GrLivArea' to is done to handle skewness.


Then we examine how SalePrice varies with Continous numeric features in dataset.



Most of the features above seems to have a good relation with SalePrice but there are some outliers present which need to be treated

After checking Discrete numeric variables with respect to salesprice in the dataset we can drop MSSubClass, YrSold & MoSold as they have no impact on SalePrice.

After all the columns are dropped from the dataset which have very less variance, then we then examine the outliers.



There are outliers in the dataset which are treated by taking the quantile of those variables.

OpenPorchSF, GarageArea, TotalBsmtSF, BsmtUnfSF, WoodDeckSF and BsmtFinSF1.

The outliers are handled by the quantile measurement of position. These numbers indicate where a specified proportion of data lies. Median is the middle position. Similarly, 25% of the data have values less than the 1<sup>st</sup> quartile whereas the 75% of data have values less than 3<sup>rd</sup> quartile.

After the outliers are handled, feature engineering on the test dataset is done where the steps like removing the id column, dividing the features into categorical and numerical data and then further dividing the numerical data into discrete and continuous features is done. Then the datetime features are handled in the same way and handling the skewness in the numeric variables is done.

Then a master dataset is created with the concatenation of training dataset and test dataset to apply encoding for the categorical variables. Encoding will basically convert the categorical variables to numeric variables. First the ordinal variables are treated first by applying integer encoding which is used to convert ordinal variables to numerical data.

```
master['ExterQual'] = master['ExterQual'].map({'Ex':5,'Gd':4,'TA':3,'Fa':2,'Po':1,'None':0})
master['BsmtQual'] = master['BsmtQual'].map({'Ex':5,'Gd':4,'TA':3,'Fa':2,'Po':1,'None':0})
master['BsmtExposure'] = master['BsmtExposure'].map({'Gd':4,'Av':3,'Mn':2,'No':1,'None':0})
master['BsmtFinType1'] = master['BsmtFinType1'].map({'GLQ':6,'ALQ':5,'BLQ':4,'Rec':3,'LwQ':2,'Unf':1,'None':0})
master['HeatingQC'] = master['HeatingQC'].map({'Ex':5,'Gd':4,'TA':3,'Fa':2,'Po':1,'None':0})
master['KitchenQual'] = master['KitchenQual'].map({'Ex':5,'Gd':4,'TA':3,'Fa':2,'Po':1,'None':0})
master['GarageFinish'] = master['GarageFinish'].map({'Fin':3,'RFn':2,'Unf':1,'None':0})
master['FireplaceQu'] = master['FireplaceQu'].map({'Ex':5,'Gd':4,'TA':3,'Fa':2,'Po':1,'None':0})
```

Then for the remaining categorical variables are encoded. Dummy Variable Encoding which is almost same as One-Hot Encoding, with one less column. There is some redundancy in One-Hot encoding. For instance, a person is either male or female. So we only need to use one of these two dummy-coded variables as a predictor. More generally, the number of dummy-coded variables needed is one less than the number of possible values, which is K-1. In statistics, this is called a dummy encoding variable, or dummy variable.

By default, the get\_dummies() does not do dummy encoding, but One-Hot encoding. To produce an actual dummy encoding from a DataFrame, we need to pass drop\_first=True. [4]

Since all the necessary operations on the train and test datasets are performed, then the dataset is again separated into train and test subset. The testing subset is to build the model. The testing subset is to use the model on unknown data for evaluation of the performance of the selected model.

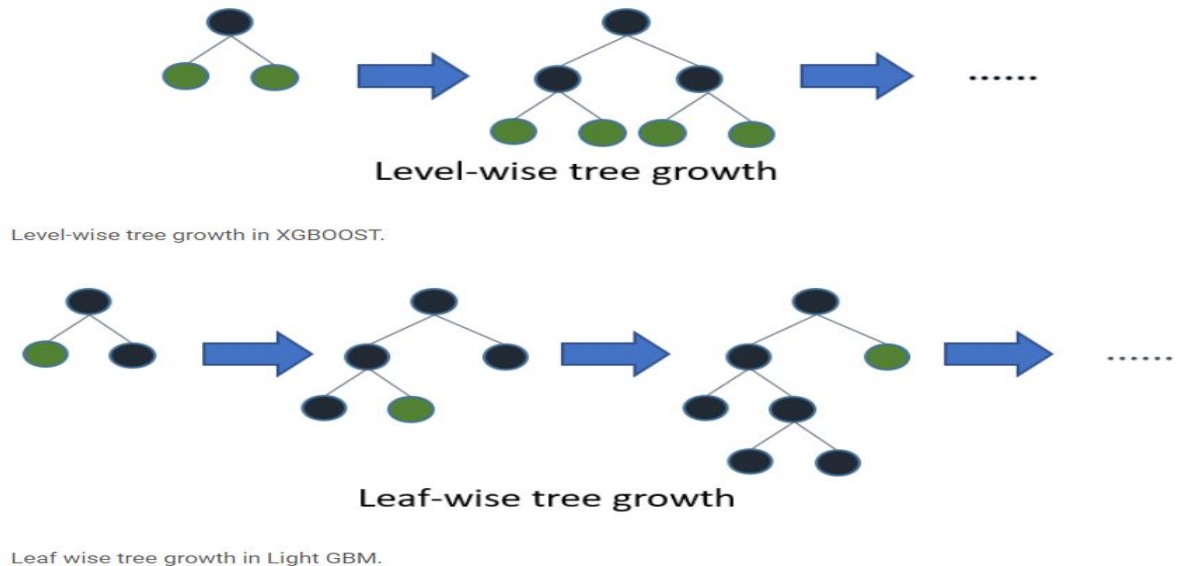
Then the `train_test_split` function is used to divide the train set into train and test data in the ratio of 70:30.

Once the train data has been split into test and train subset, a model is applied on the dataset.

First XGBoost is applied. XGBoost stands for "Extreme Gradient Boosting" and is an application of gradient boosting trees algorithm. The XGBoost is a popular SVM learning model with features like speed, parallelization, and performance. Booster stands for the regression booster type. In the model applied `gbtree` was used as this is a tree-based booster which gives better result compared to a linear model. Learning rate stands for how quick the model learns. The slower the learning rate, more better and generalized predictions are going to be and hence, 0.1 is the learning rate. Then `max_depth` will be the maximum depth of the tree which is 3 in our model. Then `N_estimators` is the count of trees generated to get the output from the model. Number need to be average as we can face overfitting or underfitting of data. `Reg_alpha` and `reg_lambda` these are the regularization values. `num_leaves` will be the maximum number of leaves for weaker model.

Light GBM is a fast, high-performance gradient boosting framework based on decision tree algorithm, used for ranking, classification. It is a decision tree algo as well, which splits the tree leaf wise with the best available fit. So when growing on the same leaf in Light GBM, it reduces more loss than the level-wise algorithm which ultimately results in higher accuracy which isn't achieved by existing boosting algorithms that's why it has the name Light.

Before is a diagrammatic representation by the makers of the Light GBM to explain the difference clearly.



## Analysis:

The coefficient of determination, which is  $r^2$  or  $R^2$  also known as R squared is applied to predict the score of the models. It is the proportion of the variance of dependent features from the target features.

The R squared prediction from XGB model is 86.62%.

The R squared prediction from LGBM model is 87.04%.

The output is saved as a csv file and before the final output is printed, the exponential values of XGB and LGBM is passed as to handle the skewness logarithmic evaluation of data was done.

In this experiment, we used 50% of XGB and 50% of LGBM values.

The kaggle score for the output file was 0.13424 ranking 2041 on the leader board.

## References

- [1] <https://www.kaggle.com/c/house-prices-advanced-regression-techniques>
- [2] <https://support.minitab.com/en-us/minitab-express/1/help-and-how-to/modeling-statistics/regression/supporting-topics/basics/what-are-categorical-discrete-and-continuous-variables/>.
- [3] <https://towardsdatascience.com/skewed-data-a-problem-to-your-statistical-model-9a6b5bb74e37#:~:text=A%20data%20is%20called%20as%20skewed%20when%20curve.For%20example%2C%20below%20is%20the%20Height%20Distribution%20graph.>
- [4] <https://towardsdatascience.com/what-is-one-hot-encoding-and-how-to-use-pandas-get-dummies-function-922eb9bd4970>
- <https://www.thoughtco.com/what-is-a-quantile-3126239>
- [https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.train\\_test\\_split.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html)
- <https://www.analyticsvidhya.com/blog/2017/06/which-algorithm-takes-the-crown-light-gbm-vs-xgboost/>
- <https://harrymoreno.com/2018/10/15/gradient-boosting-decisions-trees-xgboost-vs-lightgbm.html>